

MainRun Report

MainRun Report

Baseline

- Best validation loss (7 epochs): 1.7533 - Source run: mainrun/runs/latest - Figures (latest snapshot):

Notes

- Curves show steady training loss decrease with periodic validation spikes aligned with evals. - Validation loss converges near 1.75 by epoch 7, meeting the baseline target. - LR cosine schedule decays smoothly; throughput is stable.

baseline_20251002 — Baseline run (7 epochs) - Change: Initial logging + report pipeline - Rationale: Establish reference metrics and artifacts - Settings: `optimizer=SGD, lr=6e-3, batch=64` - Best val loss: **1.7533**

Figures:

adamw_warmup_01 — AdamW + warmup-cosine LR floor - Change: Switch optimizer to AdamW; add warmup then cosine (10% floor) - Rationale: Stabilize early steps and sustain learning late within 7 epochs - Settings: `betas=(0.9,0.95), wd=as_before, warmup≈10%/≤1000, lr_floor=0.1lr` - Best val loss: **1.4452**

Figures:

adamw_warmup_01 — AdamW + warmup-cosine LR floor - Change: Switch optimizer to AdamW; add warmup then cosine (10% floor) - Rationale: Stabilize early steps and sustain learning late within 7 epochs - Settings: `betas=(0.9,0.95), wd=as_before, warmup≈10%/≤1000, lr_floor=0.1*lr` - Best val loss: **1.4452**

Figures:





















